

SQL

- DML (Data Manipulation Language)
 - Query language
 - Update language
 - Embedded SQL(SQL/J) and ODBC(JDBV)
 - necessary for application development
- DDL (Data Definition Language)
 - defines schema for relations
 - creates(modifies/destroys) database objects
- DCL (Data Control Language)
 - access control

SQL Data Types

- Values of attributes in SQL:
 - `integer` integer (32 bit)
 - `smallint` integer(16 bit)
 - `decimal(m,n)` fixed decimal
 - `float` IEEE float (32 bit)
 - `char(n)` character string (length n)
 - `varchar(n)` variable length string(at most n)
 - `date` year/month/day
 - `time` hh:mm:ss.ss

- SQL is NOT case sensitive

```
SELECT DISTINCT <results>
FROM           <tables>
WHERE          <condition>
```

$\{ \langle \text{results} \rangle \mid \exists \langle \text{unused} \rangle. (\wedge \langle \text{tables} \rangle) \wedge \langle \text{condition} \rangle \}$

Variables vs. Attributes

- Relational Calculus uses *positional* notation, i.e.
 - `EMP(x,y,z)` is true whenever the x, y, and z components of an answer can be found as a tuple in the instance of EMP\
 - no need for attribute names
 - inconvenient for relations with high arity
- SQL uses **corelations** (tuple variables) and attribute names to assign default variable names to components of tuples:

- $R[AS]p$ in SQL stands for $R(p.a_1, \dots, p.a_n)$ in RC where a_1, \dots, a_k are the attribute names declared for R.

Naming Attributes in the Results

- Results of queries \iff Tables
- What are the names of attributes in the result of a `SELECT` clause?
 - A single attribute: inherits the name
 - An expression: implementation dependent
- We can — and should — **explicitly** name the resulting attributes:
 - \implies `"<expr> AS <id>"` where `<id>` is the new name

Boolean Connectives

- Atomic conditions can be combined using **boolean connectives**:
 - `AND` (conjunction)
 - `OR` (disjunction)
 - `NOT` (negation)

Set Operations at Glance

\implies we can apply **set operations** on them:

- set union : $Q_1 \textbf{UNION} Q_2$
 - the set of tuples in Q_1 or in Q_2
 - used to express "or"
- set difference : $Q_1 \textbf{EXCEPT} Q_2$
 - the set of tuples in Q_1 but not in Q_2
 - used to express "and not"
- set intersection: $Q_1 \textbf{INTERSECT} Q_2$
 - the set of tuples in both Q_1 and Q_2
 - used to express "and" (redundant, rarely used).

Q_1 and Q_2 must have **union-compatible** signatures:

\implies same number and types of attributes

Nesting of Queries

- We can use SELECT Blocks (and other Set operations) as arguments of Set operations.
 - What if we need to use a Set Operation inside of a SELECT Block?
 - we can use distributive laws
- $\implies (A \vee B) \wedge C \equiv (A \wedge C) \vee (B \wedge C)$
- \implies after **very** cumbersome(笨重)

- nest set operation inside a select block
 \Rightarrow common table expressions

Naming(Sub-)queries

- Syntax:

```
WITH fool [<opt-schema-1>]
  AS (<query-1-goes-here>),

  foon [<opt-schema-n>]
  AS (<quesy-n-goes-here>)

<query-that-uses-foot-...-foon-as-table-names>
```

The From clause revisited

- SQL-92 allows us to **inline** queries in the `FROM` clause:

```
FROM ... , ( <query-here> ) <id>, ...
```

\Rightarrow `<id>` stands for the result of `<query-here>`

\Rightarrow unlike for base relations, `<id>` is **mandatory**.

OR and UNION

- A common mistake:
 - use of `OR` in the `WHERE` clause instead of the `UNION` operator
- e.g.

```
SELECT title
FROM publication, book, journal
WHERE publication.pubid = book.pubid
      OR publication.pubid = journal.pubid
```

- often works; but imagine there are no `book` s...(???)

WHERE subqueries

- Additional(complex) search conditions
 \Rightarrow query-based search predicates
- Advantages
 - simplifies writing queries with negation
- Drawbacks
 - complicated semantics (especially when duplicates are involved)
 - **very** easy to make mistakes

- **VERY COMMONLY** used to formulate queries

Overview of **WHERE** Subqueries

- presence/absence of a single value in a query

Attr IN (Q)

Attr NOT IN (Q)

- relationship of a value to some/all values in a query

Attr op SOME (Q)

Attr op ALL (Q)

- emptiness/non-emptiness of a query

EXISTS (Q)

NOT EXISTS (Q)

Example

- List all authors who always publish with someone else:

```
select distinct a1.name
from author a1, author a2
where not exists(
  select *
  from publication p, wrote w1
  where p.pubid = w1.publication
    and a1.aid = w1.author
    and a2.aid not in (
      select author
      from wrote
      where publication = p.pubid
        and author <> a1.aid
    )
);
```

How do we Modify a Database?

- Naive approach:

```
DBSTART;
r1: = Q1(DB);
...
rk: = Qk(DB);
DBCOMMIT;
```

- Not an acceptable solution in practice

Incremental Updates

- Tables are large but **updates are small** \Rightarrow Incremental updates
 - insertion of a tuples (`INSERT`)
 - \Rightarrow constant tuple
 - \Rightarrow results of queries
 - deletion of tuples (`DELETE`)
 - \Rightarrow based on match of a condition
 - modification of tuples (`UPDATE`)
 - \Rightarrow allows updating "in place"
 - \Rightarrow based on match of a condition

SQL Insert

- Syntax:

```
INSERT INTO r[(a1,...,ak)]
VALUES(v1,...vk)

INSERT INTO r (Q)
```

SQL Delete

- Syntax

```
DELETE FROM r
WHERE cond
```

SQL Update

- Syntax:

```
UPDATE r
SET <update statement>
WHERE <condition>
```

- Example

```
update author
set url = 'brics.dk/~david'
where aid in (
    select aid
    from author
    where name like 'Toman%'
);
```

Support for Transactions

- The DBMS guarantees noninterference (serializability) of all data access requests to tables in a database instance
- transaction starts with first **access** of the database

⇒ until it sees:

- `COMMIT:` make changes permanent

```
SQL> commit;
```

```
Commit complete.
```

- `ROLLBACK:` discard changes

```
SQL> rollback;
```

```
Rollback complete
```

Aggregation

- Standard and very useful extension of First-Order Queries
 - Aggregate(column) functions are introduced to
 - find number of tuples in a relation
 - add values of an attribute (over the whole relation)
 - find minimal/maximal values of an attribute
 - Can apply to groups of tuples that with equal values for (selected) attributes
 - Can **NOT** be expressed in Relational Calculus
- Syntax:

```
SELECT x1,...xk, agg1,...agg1
FROM Q
GROUP BY x1,...xk
```

- Restrictions:
 - all attributes in the `SELECT` clause that are **NOT** in the scope of an aggregate function **MUST** appear in the `GROUP BY` clause.
 - `aggi` are of the form `count(y), sum(y), min(y), max(y), or avg(y)` where y is an attribute of `Q` (usually not in the `GROUP BY` clause)

Operational Reading

1. partition the input relation to groups with equal values of **grouping** attributes
2. on each of these partitions apply the **aggregate function**
3. collect the results and form the answer

- Example(count):
 - For each publication count the number of authors:

```
select publication, count(author)
from wrote
group by publication;
```

- Example(sum)
 - For each author count the number of article pages:

```
select author, sum(endpage-startpage+1) as pgs
from wrote, article
where publication = pubid
group by author;
```

- not quite correct: it doesn't list 0 pages for author 3

HAVING clause

- The **HAVING** clause is mere *SYNTACTIC SUGAR* ... and can be replaced by a nested query and a **WHERE** clause
- Example:
 - List publications with exactly one author:

```
select publication, count(author)
from wrote
group by publication
having count(author) = 1;
```

- Example:
 - For every author count the number of books and articles:

```
select distinct aid, name, count(publication)
from author, (
  ( select distinct author, publication
    from wrote, book
    where publication = pubid )
  union all
  ( select distinct author, publication
    from wrote, article
    where publication = pubid ) ) ba
where aid = author
group by name, aid;
```

Summary

- SQL covered so far:
 1. Simple SELECT BLOCK
 2. Set operations (**UNION**, **EXCEPT**, **INTERSECT**)
 3. Formulation of complex queries, nesting of queries, and views

4. Updating Data

5. Aggregation

- this covers pretty much all of the useful SQL DML
⇒ the Bad and Ugly coming next...